

Project 5

It's Extensible! Making the Medieval World a Better Place

The first part of a two-part project

Due: Monday, November 26, 2012, 11:59 PM

Note: The Corrections & Clarifications posted on the project web site are officially part of these project specifications. You should check them at least at the beginning of each work session.

Because of this project has a fairly open design, these specifications contain only few specifics of how your code should be written, and so run a risk of being under-detailed. Be sure to start early enough to have an opportunity to ask for a clarification or correction.

Purpose

This project is a extension of Project 4. You will need a working version of Project 4 as a "starter." Unless specified otherwise, Project 5 will *behave* identically to Project 4. This project will provide a foundation for the final project in the course, and so is best considered as the first phase of two.

In this project you will:

- Make use of the Singleton pattern to make Project 4's Model globally accessible in a well-controlled way to allow Sim_objects to find out about each other and indirectly broadcast to Views.
- Try out smart pointers to automate memory management using the Standard C++11 smart pointers.
- Apply the Model-View-Controller pattern more completely by implementing additional kinds of views with a full "push" approach.
- Demonstrate the value of an OOP framework by adding a new derived class of Agent, observing how only the code required to support the new class needs to be added while the remaining code is unmodified. As the gurus say, *add features by adding code, not by modifying code.*
- Get additional practice designing and refactoring your own classes.

This project assignment attempts to leave as much of the design under your control as possible. Thus the specifications will be considerably less detailed than before, and will emphasize how the program should behave, and very little about how you should accomplish it. Where necessary to make sure you work with the informative design possibilities, some design constraints are specified - a few things you must or must not do in your design.

While the design is under your control, you are expected to make good use of the OOP concepts and code quality guidelines and recommendations presented thus far in the course. You are free to modify your Project 4 code as needed. However, you should not have to make big changes to the Sim_object class hierarchy - it does most of what you need. You will have to add to this hierarchy (more details below), design a new class hierarchy for View, and modify Model's services and Controller's commands to accommodate these. In other words, your Project 5 solution should be clearly based on Project 4's, although it will change some of Project 4's code. The final project will be a further extension of this project, meaning that this is actually the first part of a two-part project. Be sure to study the Evaluation section below for information on how the code quality will be assessed.

The specifications are expressed as steps in the order you should do them for maximum benefit and smoothest work in this project.

Step 1. Make Model a Singleton class.

Once you complete this step, your program should run identically as before, but now it will be possible to do the remaining steps, which require program-wide access to information and functions in Model. Be sure you follow one of the complete recipes for the Singleton pattern, and *choose wisely* which services you will have Model provide to support the remaining steps - especially in Step 4, where a new class will need to be able to determine which objects are closest to it, meaning it will need access to information held by Model.

Note: The idiom for using a Singleton is to call its `get_instance()` function whenever you need to talk to the Singleton - so do not store a pointer or reference to it. You must remove any stored pointer or reference to Model that you had in your Project 4.

Step 2. Use Smart Pointers and Make Agents Die Immediately and Remove Themselves from the World (with Model's help).

Step 2.1. Change all containers and pointer variables used to refer to Sim_objects, Structures, Agents, etc., over to smart pointers using the C++11 Standard reference-counted smart pointer classes; see the Handouts directory for a tutorial. See the Example Code directory for examples of their use.

Who uses what kind of smart pointers? First, only `shared_ptr`s will be passed as function arguments between objects - never `weak_ptr`s. The issue is what kind of smart pointers do objects store as member variables? One way to look at this is in terms of ownership, meaning that the component that is responsible for creating and destroying the object should be the "owner" and everybody else is just an "observer." Without smart pointers, the code can be very hard to get right. But smart pointers don't help much if this ownership concept means very few `shared_ptr`s and lots of `weak_ptr`s throughout the project - working with `weak_ptr`s is

very clumsy compared to `shared_ptrs`. More importantly, the main reason for using smart pointers is so that we don't have to complicate our program design by trying to impose a "pure" idea about who "owns" what. By using shared ownership, it is easy to ensure that the object gets deleted only when nobody else is interested in it, which means that there will never be a "dangling pointer" that points to a deleted object. The only place we really need "observers" with `weak_ptrs` is to prevent smart pointer cycles between objects that refer to each other.

So here are the specifications for what kind of smart pointers each component should use:

- The Model object has containers of `shared_ptrs` to all `Sim_objects` (which includes Structures and Agents) and Views. When Model is destroyed, all of the `shared_ptrs` to `Sim_objects` and Views should get automatically destroyed when its containers are destroyed. We never use `delete` in the program - we simply let a `shared_ptr` go out of scope, reset it, erase it from a container, or let the whole container get destroyed. Review when member variable destructors are called to make sure you don't recode the wheel by doing work the compiler will do for you.
- Agents that point to other Agents should just "observe" them using `weak_ptrs`, which prevents any cycle problems that might appear with Agents that are attacking each other.
- At least in this version of the game, Structures don't refer to Agents or other Structures so there won't be any cycle problems involving Structures. So Agents will keep `shared_ptrs` to Structures. When the program is terminated, when an Structures gets destroyed may depend on when any Agents pointing to it get destroyed. By observing the destructor messages, you can see this happening.

Once you have switched over to smart pointers, you should not have any "raw" pointers for `Sim_object` family objects anywhere in your program, nor should you have any explicit deletes of these objects - the smart pointers should do this automatically, and containers of smart pointers will destroy the contained smart pointers when they get destroyed, eventually automatically deleting all the pointed-to objects. After making this change, your program should still function correctly, although dead agents may get deleted at a different time than before, depending on details in your code.

- *Hint:* See if your IDE allows you to simply do a global search/replace throughout all of the project source files of "Agent *" with "`shared_ptr<Agent>`", etc. This will save a lot of time, but you will still have to fix a few things. For example, because of how smart pointers are defined to help prevent coding mistakes, you will have to use the `reset()` member function instead of setting a pointer to `nullptr`.
- *Hint:* If you use `shared_ptr`, you will need to use the `enable_shared_from_this` facility when e.g. a Soldier calls `take_hit` on another Agent.

Step 2.2. Once you have the smart pointers in place, it is now possible to simplify how Agents disappear when they are killed: Project 4 avoided a dangling pointer problem by having the Agents go through a series of states to ensure that any other Agent (e.g. an attacker) had a chance to disconnect its pointers before the dying Agent was deleted. However, with smart pointers used throughout, there is no need for this process. In fact, as soon as an Agent realizes it is going to be dead, it should ask Model to remove it from its containers. This way, instead of Model having to be the grim reaper and monitor and remove dying Agents, Agents know when they are dying and simply ask the world (Model) to forget them immediately.

More specifically, Agents are now either **Alive** or **Dead**. There is no need for the **Dying** and **Disappearing** state; they, and their transitions, should be removed as well, along with the `is_disappearing()` accessor in Agent. Modify `Agent::update()` so that if an Agent is in the **Dead** state, it stays in that state. Modify `Agent::lose_health` from Project 4 so that if the Agent becomes no longer alive, set its state to **Dead**, and then, as the last statement before returning, call `Model::remove_agent()` with a pointer to itself! Ensure that Model removes and discards all pointers to the Agent from its containers. Subsequently, that Agent should not be updated any more or listed in any status command outputs. The code in Model for checking for and removing dead agents after each update is not needed any more and should be removed. Finally, note that in Controller, it will be impossible to try to command a non-**Alive** Agent because they will have been removed from the container of known agents immediately upon becoming non-**Alive**. Thus the Error test that a commanded Agent needs to be **Alive** is redundant and can be removed or turned into an assertion.

Step 2.3. There is a complication to this immediate removal from Model. If Model was holding the last smart pointer to the Agent, the Agent will get destroyed. However, perhaps a Soldier was holding a pointer to the now-dead Agent. If the now-dead Agent got deleted, following the pointer to ask about its internal state is undefined because the member variable values of a deleted object are undefined. Clearly, arrangements need to be made so that if another Agent is pointing to a now-dead object, the status of the dead object is well-defined.

Smart pointers make the solution easy. Soldier keeps a `weak_ptr` to its target. This kind of smart pointer doesn't keep the target in existence, but can be queried to see if the target object still exists, and a `shared_ptr` created to it if it does, which again will keep it in existence while Soldier attacks it. Check your code carefully to make sure you are correctly handling this situation - if you eliminate all raw pointers to `Sim_objects`, the compiler will call your attention to where you have to make an arrangement of this sort.

Step 2.4. The describe and update functions for Agent and Soldier need a bit of fixing. Since an Agent asks Model to remove it from the Model containers has soon as it is "killed", then the Agent will not be listed in any subsequent status commands. Thus we do not expect to see an Agent listed as being Dead - it will have already been removed from the list of Agents. However, a live Soldier might still be in Attacking state but with a target that is either dead or deleted. This could happen, for example, if another Soldier killed the target in a later update, but this Soldier has not yet been updated, and so is still in attack mode. In this project, a `weak_ptr` is used and can be tested to see if the object is still there or has been deleted. In either case, your `Soldier::describe()` should output that it is "Attacking dead target" rather than attempt to access the name of the target.

The update function needs only a slight modification from Project 4, in the checks for whether the target object is not alive. You can test to see if the target object no longer exists, or if it still exists but is dead. The target is considered "dead" in either case, and the pointer can then be reset.

Step 2.5. If Agents point to each other with `weak_ptr`s, then the cycle problem will be automatically handled. But check to see that this is really at work. A good demonstration is to have two Soldiers busy attacking each other when you issue a "quit" command. They should both get deleted. If you were using `shared_ptr`s between them, neither will get deleted; with `weak_ptr`s, they both will.

Step 2.6. Once you complete the above sub-steps, your program should run pretty much the same as before, but you will see differences in the status output and the order of destruction of Sim_objects, based on the details of which containers get emptied when, or which smart pointers get discarded at what time. In general, if you are doing this right, a killed Agent will get deleted sometime during an update cycle, not at the end of it - you should be able to set up a fight in which the destructor messages appear before some of the update messages. You will also see some differences in command error messages - for example, you will no longer be able to attempt to command a dead agent because its name will have been removed from Model - it will be an unknown Agent, not a dead Agent.

Observe the destructor messages to see this interesting "garbage collection" at work, and verify what happens when you add the remaining program features. Because exactly when the object gets deleted depends on the details of your code design, which might differ from the autograder's solution, a last step will be to remove the constructor/destructor messages - it doesn't make sense to try to match when they appear. But you should leave these message in as long as possible - it is an easy and fun way to watch the smart pointers at work and to be sure you really are releasing the object when you are done with it.

Step 3. Fully Apply Model-View-Controller and extend View to different kinds of views.

This step is an example of a common activity during software development. It is discovered that it would be a good idea to have variations on a capability already present. In this case, the View from Project 4 is just fine, but we want additional kinds of View as well, and the ability to have more than one View simultaneously active. In addition to modifying Model to handle any number of view objects, this requires refactoring the View class from Project 4 in some way. We want to retain the basic Model-View-Controller (MVC) logic in Project 4, but simply supply additional kinds of Views, and make them very flexible and under the control of the program's user. The overall structure will be very similar to how GUI code is normally organized, making this project a much better example of the MVC pattern and how it is used than Project 4. Thus, you will be designing a set of View classes and re-arranging Model to fully implement this pattern. This is a significant design project; be sure to give it some time.

Model-View-Controller in Action

It is important to follow the MVC pattern closely; review the lecture notes and the Project 4 discussion. Let's summarize the responsibilities and collaborations of the involved classes:

- **Controller** responds to user commands, and is responsible for creating Views and attaching them to the model, or detaching them and destroying them. Because the user specifies the kind of view he or she wants, Controller cares what the different kinds of Views are, but it is the **only** component that does. Controller also controls which view draws at what time or in what order, but again, it is the **only** component involved in this.
- **Model** keeps track of the Views currently attached, and will broadcast information to them, but *does not need to know and does not care what the different kinds of Views are* - so far as Model is concerned, there is only one kind of View. Giving Model any responsibilities that involve knowing anything about different kinds of Views, or in any way providing management functions for Views, means Model is doing some of Controller's job - it breaks the pattern!
- **Views**, following the push approach in Project 4, do not know anything about different kinds of Sim_objects; they merely display information about names and numeric values in a format that depends on the type of View. Model passes this name and numerical value information to them. Thus, Views are decoupled from Sim_objects in addition to Model and Controller.
- **Sim_objects** don't know anything at all about Views, but will notify Model to pass their data to the Views at appropriate times, taking advantage of Model's global availability as a Singleton.

The last bullet point above refers to a basic change from Project 4 in which Model "pushed" the location information to the View after each update for objects that *might* move, regardless of whether they actually moved or not. Instead, in Project 5, Model will rely on the Sim_objects to supply *notifications* of their data when they individually change state, by calling functions such as `Model::notify_location(const string& name, Point location)`. Model will then broadcast the notification to all of currently attached Views. You'll have a Model notification function for each type of information that Views might be interested in. This should be done as follows:

- Whenever an individual object changes in a relevant way, it will ask Model to notify all of the attached Views. For example, if Agent Zug moves, it will tell Model to notify the Views that the object named Zug is now at (20, 30). If Farm Rivendale updates to have more food, it likewise tells Model to notify the Views that the object named Rivendale now has 56 amount.
- If Controller attaches a new View to the model, then Model asks all objects to tell Model their current information to be broadcast to the Views, so that the new View is able come fully up to date (regardless of what kind of View it is).
- Similarly, if an object is added to Model, Model asks the object to notify Model of all its current information for broadcast to all current views.
- If an object is removed from Model (e.g. a dead Agent), then Model notifies all attached views that the object is gone.

Kinds of Views and New Commands

The Project 4 View, called a **map view** in this project, will still be available; the user can still control its origin, scale, and size. You can definitely recycle this code. There are three new kinds of views; see the posted samples to get a quick overview. They are:

A **local view** looks like the map view, except that it is smaller and is centered on a Sim_object's location. If the object is a Structure, then the "window" on the world stays in a fixed location. If the object is a moving Agent, then the "window" moves with the Agent, meaning that its X and Y axis labels change as the Agent moves. The user cannot change the size or scale of a local view - it

corresponds to a map view with a size of 9 and scale of 2, and its origin is adjusted to correspond to the location of the designated object. The local view output includes the object that the view is centered on. Unlike the map view, objects outside the view are not listed. If the object originally associated with a local view is gone, then the object no longer appears in the view, but the view remains centered at the last location until the user closes it.

Calculation specifics: You can use the Project 4 map view calculations by setting the map view origin x and y to the centered object's current location x and y coordinates - $(\text{size} / 2.) * \text{scale}$.

Note: The behavior of a local view for a removed object when a new object of the same name is created is undefined.

A **health view** and **amount view** are identical in format. The view lists the object names and either the current health of the object or the amount it is currently carrying. They are listed in alphabetical order by object name, with the numerical values shown with two places to the right of the decimal point for both health and amount. If an object does not have a health (like a Structure) or an amount (like a Soldier) it simply does not appear in the list - those objects simply never notify any views about this property.

Initially, no views are open - not even the map view. The user commands Controller to create, destroy, and show Views with the following commands:

- **open <view name>**. If the <view name> is map, health, or amounts, then the corresponding view is created and attached to Model. Otherwise, a local view for the named object is created and attached. If a view of this name already exists, an Error exception is thrown. If the name is not recognized as a view type or one of the current sim_objects, an Error exception is thrown.
- **close <view name>**. The view of that name is detached from model and destroyed. If no view of that name is present, an Error exception is thrown.
- **show**. Controller commands all of the current views to draw themselves. They appear in the order that they were opened.
- **default, size, zoom, pan**. Check first for an open map view, and throw an Error exception if the map view is not open. Then read and process the parameters of the command as in Project 4.

Notice that you can open any mixture of views, in any order, and a separate local view for each object known to model, and close any one of them at any time.

Automate memory management. Use smart pointers to Views to automate disposing of the Views. The end result is that there should be *no raw pointers to objects anywhere* in your Project 5!

Possible exception. Since using a smart pointer for a Singleton doesn't make any sense, if you implement the Singleton in terms of a Model * pointer, then this will be the only raw pointer in your code.

Design Goals and Constraints

For this to be the best exercise, your design must have a separate class for each kind of view, as opposed to different "modes" of fewer classes. However, you can have more than four classes if it helps achieve a good design. Plan to refactor your Project 4 View class.

The basic goal is to add these additional view capabilities to the program in a way that results in ease of extension - if we add additional different kinds of Views and different kinds of Sim_objects, we should have to modify little or no code to fit them in. Following the Model-View-Controller pattern (see above) is critical to a good design. You need to add the notification mechanism to Sim_objects in a way that meets the extensibility goals.

You also need to arrive at a good class design for the four different kinds of Views. A good way to tell whether you have a good design is to consider whether (1) you can add a new kind of Agent or Structure that has the same kinds of data (location, amount, health) with no change to either Model, Views, or Controller; or (2) you can add a new kind of view of the same kinds of data - e.g. one that showed the health of agents on a map along with their locations - without any change at all to Model or the Sim_objects, and trivial changes to Controller; or (3) we can add a View showing a new kind of data by making only the few and simple additional modifications required by the new type of data to Model and the Sim_objects.

To give you maximum flexibility in your class structure while keeping the autograder setup simple, all of your views declarations and definitions will be in a new pair of files Views.h, .cpp (notice the plural!). You won't be using View.h, .cpp any more - do not attempt to submit it. This setup violates the normal custom in which each class has its own .h, .cpp file pair; but if I dictated which separate file pairs you had to submit, it would come too close to dictating the design, which would make the exercise less valuable.

Suggestion: Set up the complete MVC pattern with the existing Project 4 map view. Get it working. You should now be able to add another view types- say the health view - with no hassle. If so, then do the design and refactoring work for the remaining View classes, implement them, and you're done!

Some Guidelines for Good OO Design

Review the notes on OO design. Some key concepts:

- Only leaf classes should be concrete - that is, don't have one concrete class inheriting from another concrete class. If this seems worth doing, it is because there is some shared functionality in the two classes that should be expressed in a base class they both inherit from.
- Put shared functionality as far up the inheritance tree as it makes sense to do. This means that the stuff done by one or more of the concrete classes should be moved up into a base class, but not so far up the tree that it becomes irrelevant for some of the other derived classes.

- Base class member functions that provide services to only derived classes should be protected. Keep base class member variables private - don't make them protected. If you are tempted to make the member variables protected, or you find you have to provide a full set of getter/setter functions for the member variables so that derived classes have full access to them, then something is wrong with the design.
- An intermediate base class should represent a meaningful abstraction in the domain. If a base class has only member variables and all of the real work is being done by functions in the derived classes using these variables, then this base class is not doing any real work and doesn't really represent anything - the nature of the member variables is an implementation detail, not a conceptual abstraction in the domain. In other words, putting just member variables in a base class does not count as "shared functionality being moved up the inheritance tree" or a "meaningful abstraction in the domain." Only member variables + substantial member functions count as shared functionality.

Rethink and Refactor. If you have problems with any of these guidelines, you need to rethink the design: what do the concrete classes actually have in common? Can you refactor the code (reorganize what the functions do) so that all of the shared work is done by the base class member functions, leaving the concrete classes to do only their own specialized bits? Maybe all they need to do is provide certain initial values, or override a little virtual function.

Step 4. Add an Archer Class.

The description of this new class is limited to describing just the behavior, so that you can design how it gets implemented, and how it relates to the existing Soldier class. An Archer is in many respects like a Soldier, but has more autonomy; it is both more aggressive and more cowardly than a Soldier. An Archer will automatically start shooting arrows (with "Twang!") at the closest Agent that is in range, but if attacked, will run away to the closest Structure. It isn't smart enough to know whether that will actually protect it or not - so if it gets attacked in Paduca, it doesn't know it should run away to somewhere else - Paduca is still the closest structure! More specifically, Archer's behavior is just like Soldier, including how it responds to commands, except for the following:

- **Initial values.** Archer has an attack strength of 1, a range of 6, and outputs "Twang!" (of a bowstring) when it hits its target (instead of Soldier's "Clang!").
- **update.** In general, Archer will check the target state, the target range, hit the target, and check again, and stop the attack if necessary, just like Soldier, but if as a result, the target is either dead or out of range, Archer will then look for a new target. More specifically, do the following in this order:
 - First update the Agent state and if Attacking, do the same checks and actions as specified for Soldier::update. If our previous target is now dead at this point, we will be in a Not Attacking state.
 - If we are now Not Attacking, we need to look for another target. Find the closest Agent (different from this one). If this Agent is within range (distance is less than or equal to the Archer range), start attacking it, outputting the message "I'm attacking". In case of a tie for closest distance, use the target whose name comes first in alphabetical order (this can be implemented trivially). Consider how to get this information and see if you can code it with a clever use of STL algorithms.
- **take_hit.** First call lose_health() and do nothing further if now Dead. Otherwise, if the attacker is still alive, find the closest structure, output "I'm going to run away to" that structure, and self-command move_to that structure. Unlike Soldiers, Archers do not automatically counter-attack their attacker. In case of a tie for closest distance, use the Structure whose name comes first in alphabetical order (this can be implemented trivially).
- **describe.** Identical to Soldier, except "Archer" appears instead of "Soldier."

Modify Model's constructor to initially create an Archer named Iriel at (20, 38), which is a convenient position for Zug to move closer - see what happens!

Design Goals and Constraints

Your design goal is to arrive at a new version of the Agent class hierarchy in which both Soldier and Archer are well represented. This is an exercise in applying the course guidelines for good Object-Oriented Design, and it should be done with care. The above list of key guidelines applies to this problem as well.

The point here is to demonstrate how you can add a new derived class with minimum changes to the rest of the code. Note that your new class should immediately work correctly with the new Views from Step 4. Therefore, you are not allowed to modify the public interface of Agent in this Step because this might break the rest of the program. However, as long as Agent and Peasant are unchanged, you are free to refactor Soldier, and/or add additional classes, if that will help achieve a good design of these "warrior" classes. The above description of Archer should not be taken as prescribing a class relationship between Archers and your previous Soldier class. You may want a new version of the Soldier class as part of the design. To allow you maximum flexibility but still keep a simple specification for the autograder, all of your code relevant for Archers and Soldiers must be in a new pair of files, Warriors.h, .cpp; the Soldier.h, .cpp files will not be used and must not be submitted. As with the Views, this setup violates the normal custom in which each class has its own .h, .cpp file pair; but if I dictated which separate file pairs you had to submit, it would come too close to dictating the design, which would make the exercise less valuable.

Step 5. Remove unnecessary constructor/destructor functions and messages.

You should keep these functions and/or messages in your program until the very end, to make sure that everything is being destroyed when it should be. Just remove them before submitting your project. As mentioned before, the order in which objects get destroyed will be rather different from Project 4, but every object that gets created should eventually be destroyed, at the earliest, when it is no longer needed, and at the latest, when the program is terminated.

Important: To demonstrate that you know when you need to define constructors and destructors versus when the compiler-supplied ones will work correctly, you need to not just delete or comment out the message output statements, but *you need to delete or comment out the entire constructor and destructor function declarations and definitions if they are not actually required for the program to work correctly (including deleting all allocated objects at termination)*.

Files to submit

Submit the same set of files that you did for Project 4 except as follows:

- the supplied **p5_main.cpp** (used as-is) instead of p4_main.cpp.
- **Warriors.h, .cpp** instead of Soldier.h, .cpp. All of your code for “attackers” must be in this .h/.cpp pair.
- **Views.h, .cpp** (notice the plural) instead of View.h, .cpp. All of your code for the views must be in this .h/.cpp pair

General Requirements

To practice the concepts and techniques in the course, your project must be programmed as follows:

- The program must be coded in Standard C++11; C Library functions should not be used (with the possible exception of `<cmath>`).
- You must follow the recommendations and correctly apply the concepts presented in this course. This means you must be familiar with the relevant material in Stroustrup, the posted C++ Coding Standards, and the lecture notes. You must review your code against these sources before submitting it.
- Your use of the Standard Library should be straightforward and idiomatic, along the lines presented in the books, lectures, and posted examples. As in Project 4, be sure to use the algorithms, iterators, bind, lambda, etc. if they can be applied simply and directly.
- You must use the smart pointers correctly, following the guidelines and concepts for their use.
- You must not use any switch-on-type logic, type codes, RTTI, or dynamic_casts in the program. Seek advice if you think you must use these or you think they would be a good idea.

Project Evaluation

This will be the last autograded project in the course. Because the design is more open, and you are allowed to modify public interfaces to the limited specified extent, component tests will not be performed. I will test your program's output to see whether your Project 4 classes and the new Archer and views behave according to the specifications.

In Project 6, the second phase, you will be choosing features to implement and designing how to implement them, so you will supply some design documentation and demonstration scripts along with your final source code. Project 6 will thus not be autograded, but human-graded throughout. The course schedule simply does not make it possible to complete a Project 5 design evaluation in time for you to respond to it with Project 6. Thus, you will get an autograder score only on Project 5, and then Project 6 will fully human-graded, but it will be evaluated for both the Project 5 design concepts and the Project 6 design requirements. You will have to build your Project 6 based on your version of Project 5. This arrangement is somewhat awkward, but past experience is that it works. One thing that makes it work is the following nasty warning:

Warning: I will spot-check that the required components and design concepts specified in this project appear to be present; if not, your autograder score for Project 5 will be *considerably* (!) reduced, possibly even being set to **zero**. In addition, I will spot-check for certain general code quality issues, such proper header file discipline and following namespace guidelines, similar to Project 2.

In other words, you should design and write Project 5 well so that you can then focus on the design problems in Project 6.

The code and design quality for both the Project 5 and Project 6 components will be weighed extremely heavily in the final project evaluation. Please study any negative feedback you got on Project 3, asking me for clarification or explanation, and make sure that your code is better this time. It is absolutely critical that your final version of these last two projects is the best that you can do.